


Silicon Sprint
Competition
Round 2

Yashvardhan Singh
Megh Ashok Giri
Siddharth Penumatsa



LOGISYM
team


Problem D1



NPU MICROCORE


4x4 Systolic Array
RTL Exploration

2025



Systolic Arrays

It is a grid of interconnected processing units, called **Processing Elements** (PEs). In this case, it's a 4x4 grid, meaning 16 PEs. Each PE receives data from its neighbors, performs a small computation, and then passes the data to its other neighbors in the next clock cycle.



Project Details

Key Design Constraints:

- Data Types:
 - Inputs (A, B): 8-bit signed integers.
 - Accumulator (Internal): 20-bit signed ($ACCW_ = 20$).
 - Output (C): 16-bit signed, saturated ($CW = 16$).
- Dataflow:
 - Matrix A elements stream in from the left and move right.
 - Matrix B elements stream in from the top and move down.
- Pipelining: Each PE must register its inputs, passing them to the next PE on the following clock cycle.


PE Design

Multiply-Accumulate (MAC):

- A combinatorial multiplier calculates the product: $\text{prod} = \text{a_in} * \text{b_in};$
- The accumulator register updates if data is valid: if (`do_acc`) $\text{c_acc} \leq \text{c_acc} + \text{prod};$


Control Signals:

- `clk`, `rst`: Standard clock and reset.
- `clr`: An extra signal to reset the 20-bit `c_acc` to zero, preparing for a new matrix calculation.
- `a_v_in`, `b_v_in`: Valid bits that are ANDed (`do_acc`) to ensure accumulation only happens when both inputs are valid.



PE Dataflow

- 1. **UNFLATTEN**: genvar i loops to convert the 1D `a_left_flat` vector into a 2D `a_left[i]` array
- 2. **BOUNDARY**: Injects this data into the edge of the PE mesh.
- assign `a_bus[i][0] = a_left[i];` (Left Edge)
- assign `b_bus[0][i] = b_top[i];` (Top Edge)
- 3. **ROW/COL Mesh**: A nested genvar r, c loop instantiates the 16 `pe_cell` modules.
- A's Output: `a_out(a_bus[r][c+1])` feeds the PE to the right.
- B's Output: `b_out(b_bus[r+1][c])` feeds the PE to the bottom.



PE Output Handling

1. Saturation (Preventing Overflow):

- Our internal accumulators (`c_acc`) are 20-bits wide to prevent overflow during summation.
- The output must be 16-bits.
- We use a `sat16` function to clip the 20-bit result to the 16-bit signed range ('sh7FFF to '-sh8000) before outputting.

2. Signaling (Knowing When We're Done):

- A systolic array has a fixed, deterministic latency.
- $\text{LATENCY} = K + (N-1) + (N-1) = 4 + 3 + 3 = 10 \text{ cycles.}$
- A state machine detects the `start_beat` (first valid data) and asserts the `done` signal exactly 10 cycles later, telling the testbench the data is ready.

RTL Issues

- **The Problem:** Our first simulations failed. The array was calculating the wrong values.
- **The Cause:** We were turning the valid signals off too early (after 4 cycles). We didn't realize the last data elements ($A[i][3]$ and $B[3][j]$) were still propagating through the pipeline
- **The Fix:** The valid signals must stay HIGH for the entire latency ($K+N-1+N-1 = 10$ cycles). This ensures all data is accumulated, even as it reaches the final PE.

=====

Systolic Array Matrix Multiplication

=====

Matrix A (4x4):

1	2	3	4
5	6	7	8
2	4	6	8
1	3	5	7

Matrix B (4x4):

1	0	0	1
0	1	0	1
0	0	1	1
1	1	1	2

Expected C = A × B:

5	6	7	14
13	14	15	34
10	12	14	28
8	10	12	23

=====

Computed C matrix:

=====

5	6	7	14
13	14	15	34
10	12	14	12
8	10	5	4

RTL Design Considerations

- Parameterized Design (parameter N=4):

The same code can generate a 2x2, 8x8, or 16x16 array just by changing parameters. This allows for rapid design exploration.

- Flattened I/O (a_left_flat [N*DW-1:0]):

Using packed 1D vectors for module I/O is a Verilog-2001 standard. This ensures maximum compatibility with synthesis tools (like Yosys) that have poor support for unpacked arrays in ports.

- generate Blocks for Mesh:

Using generate for... creates a regular, predictable structure that synthesis tools can optimize well.

Verification Strategy

1. Golden Model Calculation:

- The testbench first initializes A and B matrices.
- It computes the expected_C matrix using simple, behavioral for loops.

2. Stimulus Generation:

- The testbench asserts rst and clr to put the array in a known-good state.
- It then streams the skewed A and B matrices into the DUT, one "wave" per clock cycle, for K+N cycles.

3. Results Verification:

- The testbench waits for the done signal from the DUT and when high reads the 16-bit saturated values from C_flat.
- It compares every C[i][j] from the DUT against the expected_C[i][j].
- If all 16 values match, it prints "PASS". Otherwise, it prints "FAIL" and reports the mismatch.

Verification Testing

A for loop (from $t=0$ to $K+N-1$) feeds the correct data at the correct time.

- $t=0$: A[0][0] and B[0][0] are fed.
- $t=1$: A[0][1], A[1][0] and B[0][1], B[1][0] are fed.
- $t=2$: A[0][2], A[1][1], A[2][0] and B[0][2], B[1][1], B[2][0] are fed...
- This "wavefront" correctly mimics how the array would be fed data.

The testbench halts and does a wait(done). When done goes high, it compares its expected_C with the C_flat output from the DUT.

Results

- Successfully designed and verified a 4x4, pipelined, output-stationary systolic array.
- Architecture: The design is modular (pe_cell) and scalable (systolic_array_4x4 parameterized with N=4).
- Verification: The implementation was fully verified against a software-based golden model.

Matrix A (4x4):

1	2	3	4
5	6	7	8
2	4	6	8
1	3	5	7

=====

Computed C matrix:

=====

Matrix B (4x4):

1	0	0	1
0	1	0	1
0	0	1	1
1	1	1	2

5	6	7	14
13	14	15	34
10	12	14	28
8	10	12	23

Expected C = A × B:

5	6	7	14
13	14	15	34
10	12	14	28
8	10	12	23


Verification:

=====

PASS - All results match expected values!

=====


Results



1. **Initialization (at ~20ns):** The clr signal pulses high, correctly resetting the c_acc (accumulator) to 0 before the calculation begins.
2. **The Data-Valid Fix (at ~40ns):** The a_v_in and b_v_in signals go high and stay high for the entire streaming phase. This confirms our bug fix is implemented and the PEs are enabled for accumulation.



Results




1. **Pipelining (, a_v_in vs. a_v_out):** There is a clear 1-clock cycle (10ns) delay between a_v_in and a_v_out. This proves the PE is correctly registering its outputs, which is the core of a pipelined design.
2. **Completion (at ~140ns):** After the fixed latency, the done signal asserts high, signaling to the testbench that the calculation is complete and the final result is ready.


Simulation Environment

RTL to Netlist Automated flow on Windows

WHY?






PPA: The Language
of Chip Design



- ✗ **EDAPlayground** is **unreliable** (server downtime) → need local setup on Windows
- ✗ **Manual RTL testing** **not scalable** → need automated pipeline to avoid typing several terminal commands
- ✗ **No hardware metrics** (area, cells, timing) → need synthesis + analysis

Solution: Automated RTL-to-Netlist Flow – one-click RTL → Netlist → Reports

While our team was designing the RTL, we realized and built the infrastructure that validates and measures that work.



Simulation Environment

What we built

The Pipeline: Code → Compile → Simulate → Verify → Synthesize → Report

[IcarusVerilog] → [VVP Sim] → [Surfer Waveforms] → [Yosys+ABC] → [Area.rpt + SVG]
↓ ↓ ↓ ↓ ↓
Compile Run Test View Waveforms Gate Mapping Metrics

Execution: Single keystroke (Ctrl+Shift+B in VS Code) = entire flow automated via PowerShell!

Code + Syntax Highlighting: **VS Code**


Compilation and VCD generation: **IcarusVerilog g2012**


Waveform Viewer: **Surfer**

Synthesizer: **Yosys 58.0**

Netlist viewer: **NetlistsSVG via node.js**

Bonus: This flow is reusable
for any other Verilog design
on Windows that we want to
implement in the future.





Simulation Environment

Outputs:

Tech agnostic - ABC Area Report: **668 gate cells** (AND×28, NAND×40, NOR×257, XOR×24, XNOR×171, OR×10, NOT×27, ORNOT×42, ANDNOT×30, MUX×1), **4254+ transistors, 623 wires** considering generic 130nm values: **57,065 μm^2** ([reference](#))

PDK Synthesis pre-crash: **729 cells** (AND×23, NAND×53, NOR×304, XOR×14, XNOR×222, OR×8, NOT×15, ORNOT×26, ANDNOT×26)

Difference Observed: Gate distribution +8.9% cells, +18.3% NOR usage

Netlist Visualization:

See in next slide!

Testbench case:

All Passed!

Challenges we faced:

Challenge 1: Yosys Language Limitations

our original RTL used unpacked arrays (Yosys 0.58 doesn't fully support SystemVerilog 2012)

Solution: Flattened array structure, re-verified with testbench → PASSED the tb


Challenge 2: PDK Integration on Windows


Tried Google Skywater 130nm PDK → Liberty file syntax errors (lines 549, 8546, 13331)

Tried OpenROAD ASAP7 7nm PDK → Liberty file parsing crash (line 4151)

Root cause: Yosys Liberty parser has Windows-specific issues (known limitation)

Decision: Stick with generic gate mapping (reliable, portable)





THANK YOU